OSPYN

Your Digital Future

Regulatory Intelligence & Rule-Driven Compliance Platform

www.ospyn.com

Table of contents

1. Product Overview 3

2. Architecture Overview 4

3. Document Ingestion and Processing 5

4. AI-Driven Intelligence Layer 5

5. Workflow, Governance, and Events 6

6. Multi-Tenancy and Data Isolation 7

7. Technology Stack (Current Implementation) 7

8. The Road Ahead 8

1. Product Overview

The Regulatory Intelligence & Rule-Driven Compliance Platform is a multi-tenant SaaS system designed to continuously ingest regulatory content, interpret it using AI-assisted reasoning, and convert it into structured, versioned, and traceable rule knowledge. The platform addresses the inherent complexity of regulatory ecosystems, where policies and amendments are published as unstructured documents, evolve frequently, and must be consistently reflected across dependent artefacts. Regulatory authority content is maintained as shared, global knowledge, while each institution's internal documents, interpretations, and compliance posture remain strictly isolated.

A foundational design principle of the platform is that **rules, not documents, are the primary unit of intelligence**. Documents act as authoritative sources, while extracted rules represent actionable obligations, constraints, and thresholds. This approach enables precise amendment tracking, deterministic impact analysis, and full traceability back to original clauses. Each rule is extracted from authoritative clauses, normalized into a canonical structure, reviewed by human experts, and published into a controlled repository. When amendments are issued, the platform detects changes at the rule level and computes deltas against existing rules, rather than relying on document-level comparison.

A key capability of the system is its ability to assess downstream impact. Internal documents, interpretations, or derived artefacts are analyzed, mapped to regulatory rules, and continuously evaluated for alignment. When regulatory changes occur, the platform deterministically identifies affected downstream entities and initiates controlled remediation workflows.

The product is intentionally domain-agnostic and extensible. While it is commonly applied to compliance and regulatory monitoring scenarios, its architecture supports any use case where evolving authoritative content must be tracked, interpreted, versioned, and aligned with dependent knowledge assets in an explainable and auditable manner.

2. Architecture Overview

The platform follows a layered, service-oriented architecture with an event-driven core. Each layer has a clear responsibility and can scale or evolve independently.

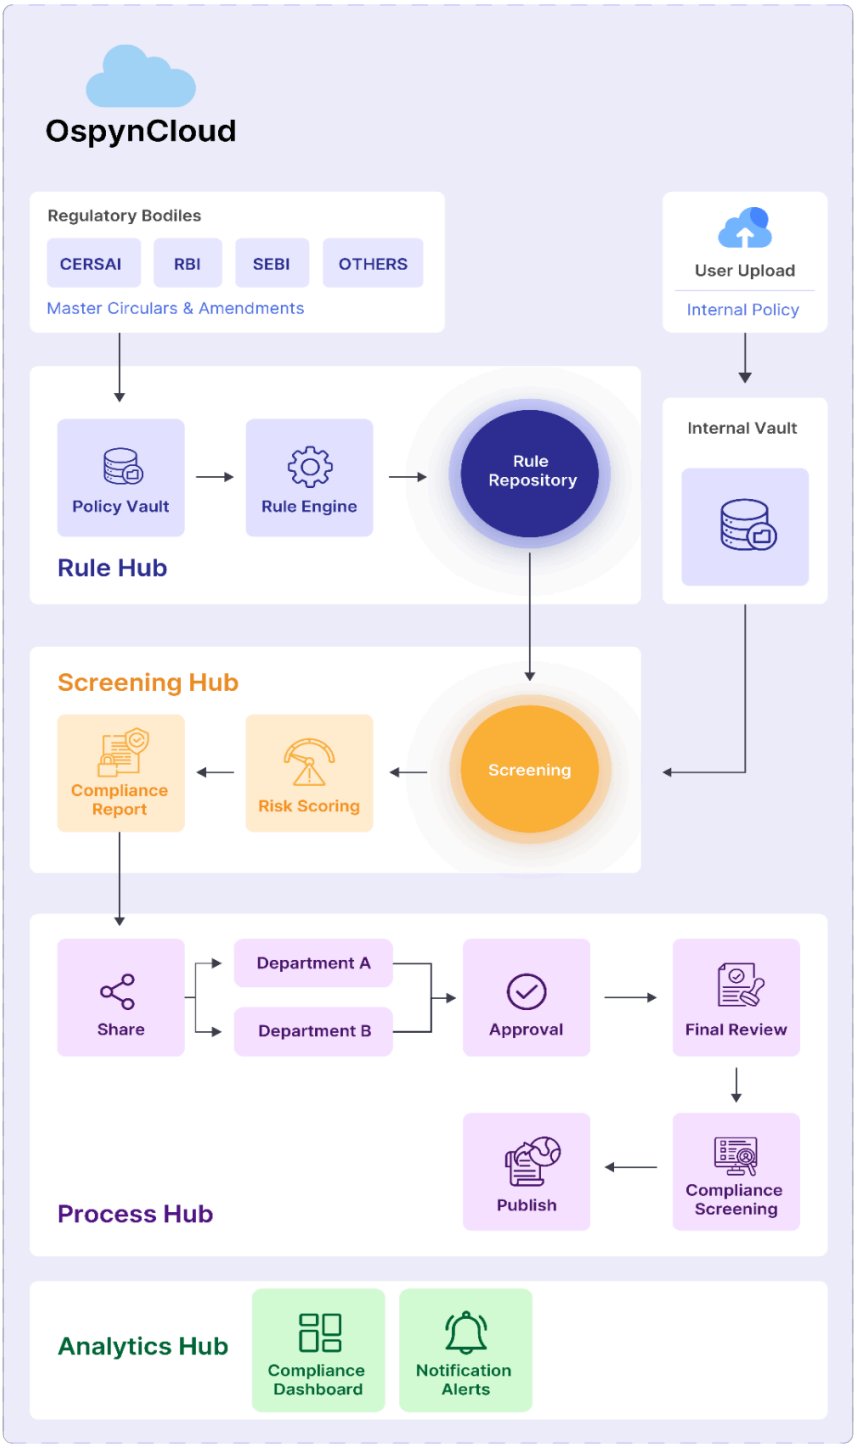

At a high level, the architecture consists of:

- A presentation layer providing web UI and APIs
- A workflow and orchestration layer governing reviews and approvals
- An AI-driven mulit-agent layer for rule intelligence
- A document ingestion and processing layer
- A hybrid knowledge repository for rules and relationships
- Immutable object storage and audit logging

3. Document Ingestion and Processing

Regulatory and institutional documents are ingested from external portals, repositories, or secure uploads. Each ingestion creates an immutable document version.

Documents are processed using layout-aware extraction to recover headings, sections, clauses, and paragraph structure. Because regulatory documents vary widely in format, the platform applies a flexible segmentation strategy:

- Structured documents (such as policies and many SOPs) are processed at clause or sub-clause level.
- Less structured documents (such as circulars or certain amendments) are processed at heading or paragraph-block level.

These segments form the base processing units that are passed into the AI intelligence layer. Large documents are processed as coordinated batches of units, allowing parallel processing while maintaining document-level consistency and traceability.

4. AI-Driven Intelligence Layer

The intelligence layer is implemented as a multi-agent system, where each agent is responsible for a clearly bounded task.

Agent Type	Responsibility
Classification Agent	Determines regulatory domain, region, and authority for a document.
Interpretation Agent	Identifies regulatory intent, obligations, constraints, thresholds, and conditions from clauses.
Rule Extraction Agent	Converts interpreted clauses into atomic, structured rules.
Amendment Comparison Agent	Compares rule versions to identify additions, modifications, and deprecations
Mapping & Impact Agent	Resolves downstream dependencies and determines impact scope.
Clarification Agent	Provides natural-language explanations grounded in citations.

All AI outputs are structured, reviewable, and governed by human-in-the-loop workflows before becoming authoritative. Human-in-the-loop checkpoints are enforced at critical stages, particularly during rule publication and amendment acceptance.

5. Workflow, Governance, and Events

The platform uses standardized, deterministic workflows to govern key activities such as rule publication, amendment acceptance, and impact resolution. These workflows ensure accountability without introducing unnecessary configuration complexity.

Inter-service communication is event-driven, allowing ingestion, analysis, and impact evaluation to scale independently. Events also serve as an additional audit trail, capturing state transitions across the system.

6. Multi-Tenancy and Data Isolation

The platform is designed as a true multi-tenant SaaS:

- Regulatory authority content, including policies and amendments, is maintained as shared, global data.
- Institution-specific artefacts, including SOPs, circulars, mappings, and workflows, are stored in tenant-isolated namespaces.
- Access controls and query enforcement ensure strict data isolation across tenants.

7. Technology Stack (Current Implementation)

The platform is implemented using a production stack that combines cloud-native services for document processing and storage with an event-driven backend and an agentic AI layer. The following components represent the current core technologies used in the platform.

Layer	Technology (Current)	Purpose
User Experience	Next.js	Provides the web UI for regulatory browsing, rule exploration, reviews, and impact dashboards.
Backend Services	Java (Quarkus)	Implements APIs, tenancy enforcement, workflows, and integration services in a cloud-native runtime.
Agentic AI Orchestration	LangGraph	Orchestrates the multi-agent execution flows as explicit graphs with inspectable state transitions.
LLM Provider	OpenAI	Provides the reasoning engine for classification, interpretation, rule extraction, delta detection, and grounded clarification.
Response Structuring	LangExtract	Enforces strict JSON outputs from LLM calls using predefined schemas for deterministic downstream processing.
Document Processing	Azure Document Intelligence	Performs OCR and layout-aware extraction to recover headings, clauses, and paragraph structure for unitization.

Layer	Technology (Current)	Purpose
Semantic Search / Vector Index	Azure AI Search	Indexes regulatory rules and tenant artefacts for similarity-based mapping, discovery, and retrieval.
Relational Store	Microsoft SQL Server (MSSQL)	Stores authoritative rule metadata, versions, approvals, mappings, workflow states, and audit references.
Messaging / Event Backbone	Apache Kafka (Confluent Cloud)	Enables decoupled ingestion, unit processing, agent execution triggers, and impact propagation via events using a managed Kafka service.
Object Storage	Azure Blob Storage	Stores immutable document originals, processed artefacts, and versioned extraction outputs with lineage references.

This stack is intentionally modular to allow the platform to evolve toward graph-based relationship storage and deeper impact traversal while preserving the existing transactional and semantic retrieval foundations.

8. The Road Ahead

The primary architectural evolution planned for the platform is the introduction of a graph-native rule repository. While the current system combines relational storage and semantic search, a graph layer will allow explicit modeling of relationships between policies, amendments, rules, and downstream artefacts. Additional roadmap items include formal rule ontologies, and automated audit narratives generated from lineage data.

Thank You

INDIA-HO & Development Centre
1st Floor, Kabani, Technopark Phase - 4, Pallippuram
Trivandrum, Kerala, India, Pin - 695 317
Ph : +91 471 667 7000

UAE
Office # 36, M Floor, Al Serkal Building
Hor Al Anz, Abu hail, Dubai , United
Arab Emirates. PO BOX: 118 457
Ph: +971 50 426 8283
www.ospyn.com | marcom@ospyn.com

USA
Ospyn Technologies, Inc.
26791 Aliso Creek Rd 1121,
Aliso Viejo, CA 92656